

OS HW#2

109370210 工管四乙 黃鈺凱

Q 4.27

The main function creates the new thread using `pthread_create` and waits for it to finish using `pthread_join`. Once the child thread finishes, the main thread outputs the generated sequence.

Note that the `fib_sequence` and `fib_count` variables are shared by both threads and therefore must be protected from race conditions. However, this program does not use mutexes or other synchronization primitives to ensure thread safety, so it may produce incorrect results when run with multiple threads. To make this program thread-safe, the `fib_sequence` and `fib_count` variables should be protected using mutexes or other synchronization primitives.

To compile the program, command:

```
gcc -pthread -o Fibonacci Fibonacci.c
```

This command will compile the `Fibonacci.c` source file and link it with the Pthread library (`-pthread` option) to produce an executable named `Fibonacci.c`.

After compilation, you can run the program with a command-line argument specifying the number of Fibonacci numbers to generate, like this:

```
huangyukai@huangyukaideMacBook-Pro programming % gcc Fibonacci.c -o Fiboacci
huangyukai@huangyukaideMacBook-Pro programming % ./Fiboacci
Usage: ./fibonacci <number of fibonacci numbers>
huangyukai@huangyukaideMacBook-Pro programming % ./Fiboacci 10
Fibonacci sequence: 0 1 1 2 3 5 8 13 21 34
huangyukai@huangyukaideMacBook-Pro programming % ./Fiboacci 20
Fibonacci sequence: 0 1 1 2 3 5 8 13 21 34 55 89 144 233 377 610 987 1597 2584 4181
huangyukai@huangyukaideMacBook-Pro programming %
```

4.24 (Optional)

Version 1

1. Program Notes:

This program will generate a series of random points based on user input. These points will be (x, y) coordinates that fall in the Cartesian coordinates in a square which has a circle inscribed in it. Based on how many points lands in the circle it generates pi.

2. To start the program:

To run the assignment, go to the terminal and find the directory that contains the compiled version of the file. To run the program, type:

```
./MonteCarlo.o
```

Then enter the desired number of points when prompted.

```
huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % gcc MonteCarlo.c -o MonteCarlo
huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % ./MonteCarlo

MultiThreaded Monte Carlo
Please enter a positive number for the amount of points you would like to generate: 100

The approximate value of pi for the desired amount of points (100) is: 3.160000

huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % ./MonteCarlo

MultiThreaded Monte Carlo
Please enter a positive number for the amount of points you would like to generate: 10000

The approximate value of pi for the desired amount of points (10000) is: 3.164400

huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % ./MonteCarlo

MultiThreaded Monte Carlo
Please enter a positive number for the amount of points you would like to generate: 1000000

The approximate value of pi for the desired amount of points (1000000) is: 3.142168
```

Version 2

In this implementation, I use four threads to generate a total of 10 million random points, if you want to change the random points just modify the (NUM_POINTS), with each thread generating 2.5 million points. The count variable is protected by a mutex to prevent data races and ensure thread safety.

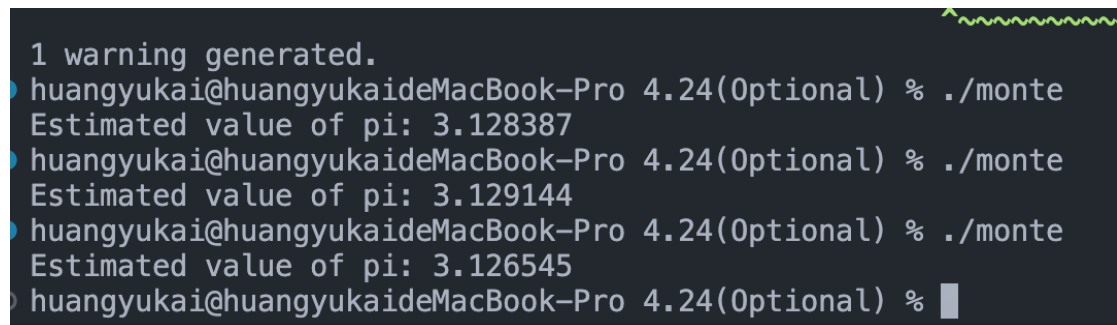
When a thread finishes, it acquires the lock on the mutex, updates the global count variable with its local count, and then releases the lock. After all threads have finished, the main thread calculates the estimated value of pi using the formula, and outputs the result to the console.

Note that the random number generator is initialized with a seed that is unique to each thread, based on the current time and the thread ID. This ensures that each thread generates a different sequence of random numbers. To create a compile file, command:

```
gcc monte.c -o monte.o -lm
```

And then, type:

```
./monte.o
```



```
1 warning generated.
huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % ./monte
Estimated value of pi: 3.128387
huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % ./monte
Estimated value of pi: 3.129144
huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % ./monte
Estimated value of pi: 3.126545
huangyukai@huangyukaideMacBook-Pro 4.24(Optional) % █
```

6.33

(a) Data involved in the race condition:

The shared resource `available_resources` is involved in the race condition. Multiple processes can access and modify this resource concurrently, leading to inconsistencies if proper synchronization mechanisms are not in place.

(b) Location of the race condition in the code:

The race condition occurs in the `decrease_count()` function where `available_resources` is decremented without any synchronization mechanism. If multiple processes try to decrement `available_resources` simultaneously, it can lead to incorrect results or data corruption.

(c) Fix the code:

The code was been fixed which show in `./6.33/count.c`.

- ◆ A mutex lock `resource_lock` is initialized and used to protect access to `available_resources`.
- ◆ The `pthread_mutex_lock()` function is used to acquire the lock before accessing `available_resources`, ensuring that only one thread can modify it at a time.
- ◆ The `pthread_mutex_unlock()` function is used to release the lock after the critical section is executed.

Programming Projects

CH4

The merge sort algorithm will have three parameters: a starting pointer, a length, and a thread-depth. For our purposes the thread depth will be N in a situation where we are using at-most $2N-1$ threads. (More on that later, but trust me, it makes it easier to do the math this way).

- If the thread depth has reached zero OR the sequence length is below a minimum threshold *we set), do not setup and run a new thread. Just recurse into our function again.
- Otherwise, split the partition. Setup a structure that holds a partition definition (which for us will be a starting point and a length as well as the thread depth which will be $N/2$), launch a thread with that parameter block, then do NOT launch another thread. instead use the current thread to recurse into `merge_sort_mt()` for the "other" half.
- Once the current thread returns from its recursion it must wait on the other thread via a join. once that is done both partitions will be done, and they can be merged using your trivial merge algorithm.

```
gary@gary-myubuntu: ~/Desktop/HW2/Ch4_proj_2$ ./MultiThreadSorting.o
542 483 199 994 981 121 874 1004
564 843 918 502 445 171 467 307
139 17 928 894 502 263 166 8
930 598 535 885 251 263 287 793
746 486 763 703 607 614 684 147
433 578 649 878 750 92 161 889
109 65 760 612 328 926 620 234
500 132 95 752 395 382 521 117
868 261 821 451 875 481 598 284
35 223 139 785 315 300 651 425
366 387 13 694 289 633 929 790
765 0 518 136 383 15 254 227
276 51 679 127 532 253 412 567
477 551 329 792 851 980 193 193
343 206 888 632 840 793 398 581
793 916 718 152 932 972 380 184
1023 35 312 531 288 724 74 765
251 403 534 78 359 727 272 702
```

```
486 978 924 746 364 158 571 107
171 174 725 280 546 955 5 496
697 977 718 391 641 971 986 949
494 349 632 585 433 803 137 919

Starting subthread...
Starting subthread...
Starting subthread...
Starting subthread...
Starting subthread...
Starting subthread...
Finished subthread.
Finished subthread.
Starting subthread...
Finished subthread.
Finished subthread.
Finished subthread.
Finished subthread.
Finished subthread.

0 0 0 0 1 1 1 2
2 2 3 5 5 5 5 6
6 6 8 8 8 9 9 10
10 12 12 12 13 13 13 13
```

949	950	950	952	953	954	954	955
955	955	956	956	957	957	958	958
959	961	961	962	962	963	963	964
964	964	965	966	966	967	968	970
970	971	971	971	971	971	972	972
972	973	973	973	973	973	973	973
974	974	974	974	975	975	976	977
977	978	978	978	978	978	979	979
979	980	980	980	981	981	981	981
982	982	982	983	984	984	985	985
985	985	986	986	986	986	986	986
986	986	987	987	988	988	988	988
989	989	990	991	991	991	992	992
992	993	993	993	993	994	994	994
995	997	999	1000	1000	1001	1001	1001
1002	1002	1003	1003	1004	1004	1005	1005
1005	1005	1006	1006	1007	1008	1009	1009
1009	1011	1012	1012	1012	1013	1013	1013
1014	1014	1014	1015	1015	1015	1016	1017
1017	1018	1019	1020	1020	1020	1021	1021
1021	1022	1023	1023	1023	1023	1023	1023

The most important part of this is the limiters that keep us from going thread wild, which is easy to accidentally do with recursive threaded algorithms, and the join of the threads before merging their content with the other half of the partition, which we sorted on our thread, and may also have done the same thing.

CH5

The code first creates copies of the original process list for each scheduling algorithm. This is done to ensure that each algorithm operates on the same set of processes, allowing for a fair comparison of their performance.

After the process lists are copied, the program runs each scheduling algorithm in turn, printing the results to the console. The FCFS, SJF, and Priority Scheduling algorithms are run without any additional user input.

For the Round Robin Scheduling algorithm, the program prompts the user to enter a time quantum. The time quantum is a parameter of the Round Robin algorithm that determines how long each process is allowed to run before it is interrupted and moved to the back of the queue. After the user enters the time quantum, the Round Robin algorithm is run and its results are printed to the console.

If we enter the data like this

```
huangyukai@huangyukaideMacBook-Pro CH5 % ./process
Enter the number of processes: 5
Enter arrival time, burst time, and priority for process 1: 0 3 2
Enter arrival time, burst time, and priority for process 2: 2 6 1
Enter arrival time, burst time, and priority for process 3: 4 4 3
Enter arrival time, burst time, and priority for process 4: 6 5 2
Enter arrival time, burst time, and priority for process 5: 8 2 4
```

◆ FCFS Scheduling:

```
FCFS Scheduling:
Process Arrival Time    Burst Time    Waiting Time    Turnaround Time
1      0              3              0                3
2      2              6              3                9
3      4              4              9               13
4      6              5             13               18
5      8              2             18               20
Average Waiting Time: 8.60
Average Turnaround Time: 12.60
```

◆ SJF Scheduling:

SJF Scheduling:

Process	Arrival Time	Burst Time	Waiting Time	Turnaround Time
1	0	3	0	3
2	2	6	1	7
5	8	2	1	3
3	4	4	7	11
4	6	5	9	14

Average Waiting Time: 3.60
Average Turnaround Time: 7.60

◆ Priority Scheduling:

Priority Scheduling:

Process	Arrival Time	Burst Time	Priority	Waiting Time	Turnaround Time
P2	2	6	1	0	6
P1	0	3	2	6	9
P4	6	5	2	3	8
P3	4	4	3	10	14
P5	8	2	4	10	12

Average Waiting Time: 5.80
Average Turnaround Time: 9.80

◆ Round Robin Scheduling (time quantum = 4)

Enter time quantum: 4

Round Robin Scheduling:

Process	Arrival Time	Burst Time	Waiting Time	Turnaround Time
P1	0	3	0	3
P3	4	4	3	7
P5	8	2	7	9
P2	2	6	11	17
P4	6	5	9	14

Average Waiting Time: 6.00
Average Turnaround Time: 10.00

◆ Priority with Round Robin Scheduling:

Priority with Round Robin Scheduling:

Process	Arrival Time	Burst Time	Priority	Waiting Time	Turnaround Time
1	0	3	2	0	3
2	2	6	1	1	7
4	6	5	2	3	8
3	4	4	3	10	14
5	8	2	4	10	12

Average Waiting Time: 4.80
Average Turnaround Time: 8.80

Summary:

Team members

工管四乙 黃鈺凱 109370210 50%

資工三 林兆玄 110590021 50%